

Andrew R. Mains

510-730-9376 | amains12@gmail.com | github.com/andrewmains12

EDUCATION

UC Berkeley

B.A. Computer Science & Cognitive Science, GPA: 3.75

Berkeley, CA

December, 2012

SKILLS AND TECHNOLOGIES

Languages: Java, Python, Scala, C/C++, Ruby, Clojure, SQL, Scheme, R

Technologies: Hadoop, Hive, HBase, HCatalog, Saltstack, ANTLR, Bison/Flex, MySQL, Django, Jersey, Rails

EXPERIENCE

Uber

Software Engineer II, Marketplace Continuity

San Francisco, CA

January 2016-present

Flipr: Dynamic Configuration

- Fullstack development on Uber's dynamic configuration service

Upsight (formerly Kontagent)

Senior Software Engineer, Data Team

San Francisco, CA

August 2014-December 2015

Software Engineer, Data Team

February 2013-August 2014

Tesseract: OLAP over HBase (Java, HBase, Jersey)

- Implemented a query system for multidimensional data on top of HBase with support for rollups, grouping and filtering
- Implemented an HBase coprocessor to distribute query execution across region servers
- Architected and implemented a REST API on top of the underlying query layer

Datamine Hardening (Django, Hive, HCatalog)

- Led a team of 4 in redesigning and implementing an asynchronous REST API for Kontagent's multitenant hive service, focusing on reliability
- Patched Apache's WebHCat to support asynchronous submission of jobs to Hadoop's JobTracker

Oracle

Software Engineering Intern

Redwood Shores, CA

Summer 2012

SAS to R Compiler (ANTLR, Java)

- Architected and implemented a translator from SAS to R from the ground up
- Created a static analyzer for SAS in order to generate optimized R code.

Kontagent

Software Engineering Intern

San Francisco, CA

Summer 2011

Hive Interface (Python, Hive)

- Wrote the server side of a REST interface for Hive, using the Thrift API provided by Cloudera's Beeswax
 - Wrote a UDF to calculate session times for users in games
-

OPEN SOURCE CONTRIBUTIONS

Mapreduce over multiple HBase snapshots (HBASE-13356, HBase 1.2.0)

- Implemented the ability to run mapreduce jobs over multiple scans on HBase snapshots
- Allowed for more fine grained push down of predicates to HBase, and thus better performance

Hive predicate pushdown to multiple HBase scans (HIVE-7805)

- Modified Hive's HBase integration to push query predicates down to multiple scans
 - Significantly reduced the amount of data scanned for hive queries on Upsight's raw data store
-

SELECTED COURSES

CS162: Operating Systems

CS164: Compilers

CS170: Algorithms

CS188: Artificial Intelligence

CS186: Databases

CS169: Software Engineering